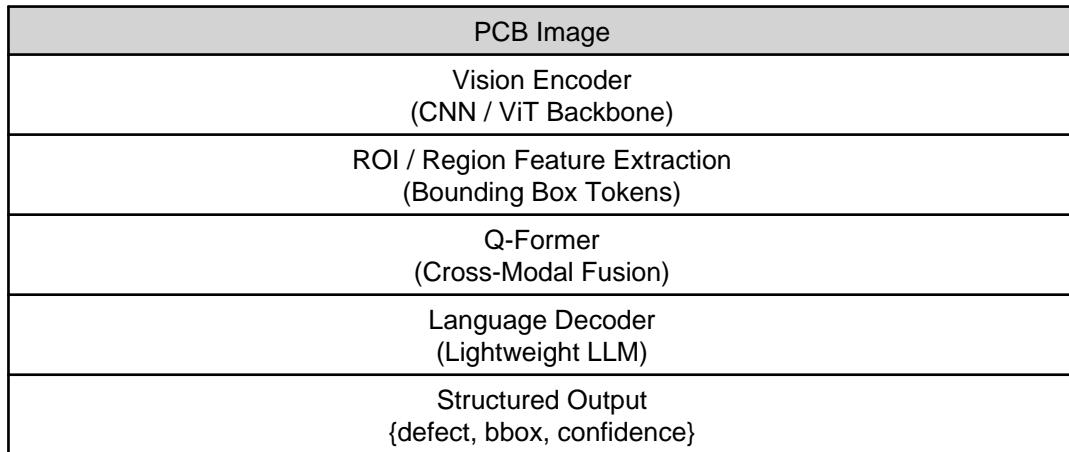


Assignment 3: Custom Vision–Language Model (VLM) Design for Industrial Quality Inspection

Candidate: Bashwaraj Sonkawade

System Architecture Overview



The architecture follows a modular BLIP-2-style design where the vision encoder extracts PCB-specific visual features. Region-level embeddings are created using bounding box information, which are fused with language tokens via a Q-Former module. A lightweight language decoder produces structured, grounded responses suitable for industrial inspection.

Key Design Principles

- Region-grounded visual representations
- Structured language generation (no free-form hallucination)
- Lightweight and offline deployable
- Sub-2 second inference latency